

**NANYANG
TECHNOLOGICAL
UNIVERSITY**

**NANYANG TECHNOLOGICAL UNIVERSITY,
SCHOOL OF COMPUTER SCIENCE AND ENGINEERING**

CZ4031 Database System Principles

Project 1

Group 41

Group Members:

Ong Kwang Wee (U2021723J)

Wong Jing Yen (U2020809H)

Kenny Tan Junrong (U1920056G)

Manimaran Manoj Kumar (U2021005E)

Alloysius Lim Zong Hong (U1922302F)

| | |
|---|-----------|
| Contribution of Team Members | 3 |
| Introduction | 4 |
| Description | 4 |
| Project Overview | 4 |
| Program Run Instructions | 4 |
| Design of Storage Components | 5 |
| Records | 5 |
| Blocks | 5 |
| Disk Storage | 7 |
| Disk Operations | 8 |
| Design of B+ Tree Components | 8 |
| Data Structure of a node | 8 |
| Maximum number of keys in a node | 8 |
| Building of B+ Tree (Insert, Search & Delete) | 9 |
| Experiment Results | 20 |
| Experiment 1 Results | 20 |
| Experiment 2 Results | 21 |
| Experiment 3 Results | 22 |
| Experiment 4 Results | 28 |
| Experiment 5 Results | 35 |

Contribution of Team Members

| Task | Name |
|------------------------------------|--|
| Storage Implementation | Ong Kwang Wee Kenny Tan Junrong |
| Base of B+ Tree Implementation | Wong Jing Yen |
| B+ Tree Implementation (Insertion) | Wong Jing Yen Kenny Tan Junrong Alloysius Lim Zong Hong Manimaran Manoj Kumar |
| B+ Tree Implementation (Search) | Wong Jing Yen Kenny Tan Junrong Manimaran Manoj Kumar |
| B+ Tree Implementation (Deletion) | Wong Jing Yen |
| Experiments | All members |
| Report | All members |

| Name | Contribution |
|-------------------------|--------------|
| Wong Jing Yen | 20% |
| Kenny Tan Junrong | 20% |
| Manimaran Manoj Kumar | 20% |
| Ong Kwang Wee | 20% |
| Alloysius Lim Zong Hong | 20% |

Introduction

Description

This current project focuses on the design and implementation of the two components of a database management system, mainly storage and indexing. It also demonstrates a B+ Tree implementation that supports inserting, searching and deleting operations.

Project Overview

The implementation consists of the following classes with the corresponding responsibilities:

- Main.java: The boundary class that accepts user input and invokes methods in other classes to conduct the experiments for the project
- BPlusTreeNode.java: Implementation of the B+ Tree structure through insertion and deletion of keys into nodes. It also supports searching of keys.
- Disk.java: Supports insertion and deletion of records into blocks. Also supports retrieval of data from specified records from the disk.
- Block.java: Implementation of Block structure that controls the storage of records
- Record.java: Implementation of Record structure that holds the fields (tconst, averageRating, and numVotes).
- Utils.java: To map unchanging values to respective variable constants that are used throughout the classes that aids to make the code more robust.

Program Run Instructions

This program is built and compiled to a .jar file type to ensure the ease of running the program. The steps below are to run the .jar file from Terminal or IDE.

Clone the repository from Github repository

<https://github.com/kennytaan/CZ4031-DBSP-Java> or extract the zipfile submitted. Ensure that you have any supporting Java IDE (E.g. Microsoft Visual Studios, Eclipse, IntelliJ) which may require additional installations like JDKs to open the project folder. A recommended installation guide is provided on README.md on github repository.

These are the following steps to run the program:

1. Navigate to the Main.java.file on IDE or Terminal (PATH: /CZ4031-DBSP-Java)
2. Running Main.java through the IDE or Terminal
3. Experiments 1-5 will be carried out with disk block size of 200MB and respective outputs can be seen under the Terminal after it has completed running.
4. The terminal will prompt for an user input to:
 - [1] Repeat Experiments 1-5 with disk block size of 500MB
 - [2] Exit the program
5. The program will continue to run only if "1" is entered, else it will prompt for the correct user input or exit the program.

- Experiments 1-5 will be carried out with disk block size of 500MB and respective outputs can be seen under the Terminal after it has completed running.

Design of Storage components

Records

The data.tsv file consists of 3 column values: tconst, averageRating, numVotes. Each line read from the tsv file is split into fields which are then packed into a single record as follows:

- tconst -> e.g. tt9916778 -> char[] (10characters with 2bytes each = 20bytes)
- averageRating -> e.g. 9.9 -> float (4 bytes)
- numVotes -> e.g. 4582 -> int (4 bytes)

Total record size = 28 bytes

```
byte[] tConst_b = stringToBytesUTFCustom(record.getTConst());
byte[] avgRating_b = floatToByteArray(record.getAverageRating());
byte[] numVotes_b = intToByteArray(record.getNumVotes());

ByteArrayOutputStream byteOutputStream = new ByteArrayOutputStream();
byteOutputStream.write(tConst_b);
byteOutputStream.write(avgRating_b);
byteOutputStream.write(numVotes_b);
byte[] record_b = byteOutputStream.toByteArray();
```

Figure 1: Logic to handle conversion of datatype to Byte Array

As show in Fig 1, each record attribute is being converted to a byte array to ensure consistency in data size. Thus, no additional memory is required to indicate the size of each field.

Blocks

```
public class Block {
    private byte[] data;
    private int blockOffset;
```

Figure 2: Block class

The class Block consist of 2 private declarations:

- data
 - Consist of the data bytes of a single record
- blockOffset
 - Acts as a pointer to determine the next location of the next record. The offset will be adjusted if a new block is created.




Figure 3: How records are packed into block

This project utilises the unspanned record method to pack records into blocks. This means that the record can only be stored inside a block if it can be completely stored into it. A record cannot be stored in more than one block. The reasons for choosing this method are as follows:

- Record size is smaller than block size.
- Access time of a record is much faster.
- It is a simpler way for storage implementation as compared to the spanned record method.

```
int offset;  
  
if (blocks[currentBlock] == null) {  
    blocks[currentBlock] = new Block();  
}  
  
if (blocks[currentBlock].getOffset() < (BLOCKSIZE - RECORDSIZE)) {  
    offset = blocks[currentBlock].getOffset();  
    blocks[currentBlock].addData(record_b);  
} else {  
    currentBlock++;  
    blocks[currentBlock] = new Block();  
    offset = blocks[currentBlock].getOffset();  
    blocks[currentBlock].addData(record_b);  
}
```

Figure 4: Logic to handle packing of records into a block

Fig 4 shows how we implement the part where the records can only be inserted into the block if its size is smaller than the offset. A new block will be created to store the record if the record's size is greater than the offset and cannot be inserted into the block.

Disk Storage




Figure 5: How blocks are stored in a disk

The blocks are then arranged sequentially on the disk to minimise seek and rotational delay.

```
//store blocks in an array
// to access certain block --> blocks[blockno]
private Block[] blocks;
private int currentBlock;
```

Figure 6: Block class

Similar to the diagram shown in Fig 5, Fig 6 shows how we implement the disk storage by using an array of Block objects to store the blocks and the index of the array represents the current block that we want to access. For e.g. when current block(int) = 4, blocks[4] allow us to access the 5th block in the disk.

Disk Operations

The disk is able to support several operations:

- Insert Record
- Delete Record
- Search Record
- Print Record

Design of B+ Tree Component

Data Structure of a node

```
private int[] keys = new int[NUMOFKEYS];  
private Object[] pointers = new Object[NUMOFPOINTERS];  
private int size = 0;  
private int height = 0;  
//private static int numOfDeletedMerged = 0;
```

Figure 7: Variables of a node

Every Node consists of the following variables:

- Keys
 - o Integer Array of size (Bytes) that stores the (name) value
- Pointers
 - o Object Array of size (Bytes) that stores pointers address of each child/sibling nodes or (name) value
- Size
 - o Size of the node
- Height
 - o Height of the B+ Tree

Maximum number of keys in which a Node maintains


```

/*
BLOCKSIZE
n+1 pointers (8 Bytes)
n integer keys (4 Bytes)
BLOCKSIZE = 8n + 8 + 4n
1 BLOCK can fit in (BLOCKSIZE-8)/12 keys
*/
public static final int NUMOFKEYS = ((BLOCKSIZE-8)/12);

public static final int NUMOFPOINTERS= NUMOFKEYS+1;

public static final int MINKEYS = (NUMOFKEYS+1)/2;

```

Figure 8: Functions in Utils.java

In Utils.java, functions are implemented to handle the calculation of values for respective variables. Fig 8 is how we calculate the number of keys, number of pointers and minimum number of keys in a node. As we know that the pointer size is 8bytes due to 64-bit processor and operating system, together with the key size which is 4 bytes for numVotes integer, we can know that every key-pointer pair will have a size of 12bytes in total. Hence, we are able to calculate the maximum number of keys that can be implemented in a node with regards to the block size. At the same time, we are also able to calculate the minimum number of keys in a leaf node with the formula which is the floor of $(n+1)/2$.

```

public class BTreeNode {

    public static final int MAX_KEYS = NUMOFKEYS;
    private static final int MAX_POINTERS = NUMOFPOINTERS;
    private static final int MIN_KEYS = MINKEYS;
    private static final int MIN_NON_LEAF = NUMOFKEYS/2;
    private static int root = 0;
    private static int numOfNodes = 0;
}

```

Figure 9: Variables instantiated in BTreeNode class

With the values we get from Utils.java, the maximum number of keys (MAX_KEYS) can be fetched from it, which correlates to the size of the block.

- For a block size of 200B, MAX_KEYS will be $\lfloor (200 - 8) / (4 + 8) \rfloor = 16$
- For a block size of 500B, MAX_KEYS will be of $\lfloor (500 - 8) / (4 + 8) \rfloor = 41$

Building of B+ Tree

The implementation of B+ Tree is mostly done in BTreeNode.java. It supports the 3 main operations: Insert, Search and Delete.

Insertion

```
//only for leaf node
//insert a key and sort the node
public void insertKey(int key, Object pointer){
    for(int i=0;i<MAX_KEYS;i++){
        if(key < this.keys[i] || this.keys[i] == -1){
            for (int j=this.size;j>i;j--){
                this.keys[j] = this.keys[j-1];
                this.pointers[j] = this.pointers[j-1];
            }
            this.keys[i] = key;
            this.pointers[i] = pointer;
            this.size++;
            return;
        }
    }
}
```

Figure 10: Base logic to handle insertion of keys in leaf node

```
//only for non leaf node
//insert a key with both child pointers and sort the node
public void insertKey(int key, Object pointer1, Object pointer2){
    int i;
    for(i=0;i<this.size+1;i++){
        if(key < this.keys[i] || this.keys[i] == -1){
            for (int j=this.size;j>i;j--){
                this.keys[j] = this.keys[j-1];
                this.pointers[j + 1] = this.pointers[j];
                this.pointers[j] = this.pointers[j-1];
            }
            this.keys[i] = key;
            this.pointers[i] = pointer1;
            this.pointers[i+1] = pointer2;
            this.size++;
            break;
        }
    }
}
```

Figure 11:Base logic to handle insertion of keys in non-leaf node

Fig 10 and 11 shows the logic to handle insertion of key into both leaf and non-leaf nodes and to sort them after it is being inserted. By looping through the keys in the node, the new key inserted will be compared key by key. This is to check the actual position that it should be inserted within the node.

```

public BTreeNode insertLeafNode(int key, int address){
    // if node is not full insert key
    if (this.size < MAX_KEYS){
        this.insertKey(key, address);
        return this;
    }
    // node is full
    numOfNodes++;
    BTreeNode newLeaf = new BTreeNode();
    int newNodeSize=0;
    //split the node
    for(int i=MIN_KEYS;i<MAX_KEYS; i++){
        newLeaf.keys[newNodeSize] = this.keys[MIN_KEYS];
        newLeaf.pointers[newNodeSize] = this.pointers[MIN_KEYS];
        this.deleteLeafPointer(this.pointers[MIN_KEYS]);
        newNodeSize++;
    }
    newLeaf.size = newNodeSize; //new leaf is right node

    //insert key into corr node
    if (key < newLeaf.keys[0]){
        this.insertKey(key, address);
    }
    else{
        newLeaf.insertKey(key, address);
    }
    // check this.size == newLeaf.size or this.size == newLeaf.size+1
    if (this.size < newLeaf.size){
        this.insertKey(newLeaf.keys[0], newLeaf.pointers[0]);
        newLeaf.deleteLeafPointer(newLeaf.pointers[0]);
    }
    else if(this.size> newLeaf.size+1){
        newLeaf.insertKey(this.keys[this.size-1], this.pointers[this.size-1]);
        this.deleteLeafPointer(this.pointers[this.size-1]);
    }
    // arrange last pointer
    newLeaf.pointers[MAX_KEYS] = this.pointers[MAX_KEYS];
    this.pointers[MAX_KEYS] = newLeaf;

    //return parent node
    BTreeNode parent = new BTreeNode();
    parent.insertKey(newLeaf.keys[0], this, newLeaf);
    parent.height = 1;
    if(this.root < parent.height) {
        this.root = parent.height;
        numOfNodes++;
    }
    return parent;
}

```

Figure 12: Logic to handle insertion of keys into leaf node

Fig 12 shows how we implement the insertion of keys into a leaf node. If the leaf node still has empty space for insertion, we will just insert the key into the directly insert the key into the leaf node and return the node.

However, if we know that the insertion will cause an overflow, we will split the node and insert the second half of the node into a corresponding new node. We will then compare and check whether the key should be inserted in the left node or the right node. After inserting the key into the correct node, we will create a new parent node which holds the value of the minimum key of the right node and return it.

```

//insert the key here
public BTreeNode insertNode(int key, int address){
    // if this is leaf
    if (this.height == 0){
        if(this.root == 0 && this.size == 0) numOfNodes++;
        return this.insertLeafNode(key, address);
    }
    // get child
    int childIndex = 0;
    for(int i=0;i<this.size && key >= this.keys[i]; i++){
        childIndex = i+1;
    }
    // this is a non leaf
    //insert to child
    BTreeNode child = (BTreeNode) this.pointers[childIndex];
    this.pointers[childIndex] = child.insertNode(key, address);
    // if same parent key is returned
    if (child == this.pointers[childIndex]){
        return this;
    }
    BTreeNode newParent = (BTreeNode) this.pointers[childIndex];
    // if current non leaf is not full
    if( this.size< MAX_KEYS){
        this.insertKey(newParent.keys[0], newParent.pointers[0], newParent.pointers[1]);
        return this;
    }
    // number of keys allocated to new right node
    int newNodeKeys = MAX_KEYS/2;
    int newNodePointers = newNodeKeys+1;
    // number of keys allocated to left node
    int oldNodeKeys = MAX_KEYS-newNodeKeys;
    int oldNodePointers = oldNodeKeys+1;
    int j=0;
    // copy all keys and pointers to one array
    BTreeNode[] childNodes = new BTreeNode[MAX_KEYS+2];
    for(int i=0; i<MAX_POINTERS;i++){
        if(i==childIndex){
            childNodes[j++] = (BTreeNode) newParent.pointers[0];
            childNodes[j++] = (BTreeNode) newParent.pointers[1];
            continue;
        }
        childNodes[j++] = (BTreeNode) this.pointers[i];
    }
}

```

Figure 13(i): Logic to handle insertion of keys into the B+ Tree (Part 1)

Fig 13(i) shows the first part of how we implement the insertion of keys into the B+ Tree. If we know that the tree is on the first level by checking the height, we know that it is a leaf node and call the insertion of keys into the leaf node as shown in Fig 12.

If it is not the first level, the function will retrieve the child nodes first and try to insert the keys into the child node with the same parents. The number of keys will be updated and allocated to the left node and the new right node. All the keys and pointers will then be copied into an array.

```

// initialise left non-leaf node
// which we are going to reuse the current node
this.pointers[0] = childNodes[0];
for(int i=0; i<MAX_KEYS; i++){
    if(i < oldNodeKeys){
        this.keys[i] = findMin(childNodes[i+1]);
        this.pointers[i+1] = childNodes[i+1];
    }
    else{
        this.keys[i] = -1;
        this.pointers[i+1] = null;
    }
}
this.size = oldNodeKeys;

//initialise right non-leaf node
// we are going to create a new node for this
BTreeNode newNonLeaf = new BTreeNode();
newNonLeaf.height = this.height;
newNonLeaf.size = newNodeKeys;
newNonLeaf.pointers[0] = childNodes[oldNodePointers];
for(int i=0; i<newNodeKeys; i++){
    newNonLeaf.keys[i] = findMin(childNodes[oldNodePointers+i+1]);
    newNonLeaf.pointers[i+1] = childNodes[oldNodePointers+i+1];
}

//new sibling node created add nodes
numOfNodes++;

//return parent node
BTreeNode parent = new BTreeNode();
parent.insertKey(findMin(childNodes[oldNodePointers]), this, newNonLeaf);
parent.height = this.height + 1;
// if new root is created
if (this.root < parent.height){
    this.root = parent.height;
    numOfNodes++;
}
return parent;
}

```

Figure 13(ii): Logic to handle insertion of keys into B+ Tree (Part 2)

Searching

```
public int[] search(int min, int max, boolean printNodes){
    // count number of access
    int count=0;
    if (printNodes){
        System.out.println(x: "Accessing nodes:");
    }
    ArrayList<Integer> result = new ArrayList<Integer>(Arrays.asList(-1));
    BTreeNode curNode;
```

Figure 14: Searching of leaf nodes (Declaration)

Fig 14 shows the logic to handle searching of both leaf and non-leaf nodes. Firstly, an ArrayList is declared in order to store the results and the BTreeNode is called instantiated in order to traverse down the B+ tree.

```
int ptr=-1, i;
//get the min leaf node
curNode = this;
while(curNode.height>0){
    if(printNodes){
        //increment accesses
        count++;
        System.out.printf(curNode.getContent()+"\n");
    }
    ptr =0;
    for(i=0; i< curNode.size && min >= curNode.keys[i];i++){ // scan the node for child with key
        ptr = i+1;
    }
    if (i== curNode.size) ptr = curNode.size;
    curNode = (BTreeNode) curNode.pointers[ptr];
}
```

Figure 15: Searching of leaf nodes (Finding the leaf Node)

Fig 15 shows the continuation of the logic to handle searching of both leaf and non-leaf nodes. To traverse down the B+ tree, a loop is used against the height of the tree to recursively access the minimum leaf node. To ensure that the minimum leaf node is accessed, the count of the traversing is checked against the size of the current node.

```

while(curNode!=null){ // scan leaf nodes
    if(printNodes){
        count++;
        System.out.printf(curNode.getContent());
    }
    for(i=0;i< curNode.size;i++){
        if(curNode.keys[i] >= min && curNode.keys[i] <= max) { // if within range add to list
            if (result.get(index: 0) == -1) result.remove(index: 0);
            result.add((Integer) curNode.pointers[i]);
        }
        else if(curNode.keys[i] > max || curNode.keys[i] == -1){ //if max is reached
            if(printNodes)
                System.out.println("Number of Nodes accessed: "+ count);
            return result.stream().mapToInt(num->num).toArray();
        }
    }
    curNode = (BTreeNode) curNode.pointers[MAX_KEYS]; // go to next node
    if(curNode == null){
        if (printNodes) System.out.println(x: "Reached end of database");
        else if (result.get(index: 0) == -1) System.out.println(x: "key not found!");
    }
}
System.out.println("Number of Nodes accessed: "+ count);
return result.stream().mapToInt(num->num).toArray();
}

```

Figure 16: Searching of leaf nodes (Finding the leaf Node)

Fig 16 shows the continuation of the logic to handle searching of both leaf and non-leaf nodes. While accessing the leaf nodes of the B+ tree, the key value of the leaf nodes are checked to see if it is within the range. If the values are indeed within the range of the leaf nodes, the key-value pair will be pushed into the `ArrayList<result>` and return as the output. If the values exceeds the maximum value of the leaf node, it will be redirected to access the subsequent leaf node until the value could be found within the list range. After all the leaf nodes have been accessed and if the values could not be found, `search()` will return a message to indicate the end of the database with no key found.

Deletion

```

public BTreeNode remove(int key, ArrayList<BTreeNode> parents, ArrayList<Integer> parentPointer){
    BTreeNode curNode = this;
    int i=0;
    int ptr;
    // this gives me the leaf node
    while(curNode!=null) {
        while (curNode.height != 0) {
            //this gets the child position in parent
            // we need to check the node before to see if duplicates are there
            for (i = 0; i < curNode.size && key > curNode.keys[i]; i++) {
                continue;
            }
            //remember the parent node and the position of the child node in parent
            //treat this list as a stack
            parents.add(index: 0, curNode);
            parentPointer.add(index: 0, i);

            curNode = (BTreeNode) curNode.pointers[i]; //child node with key
        }
        // find key and delete in leaf node
        boolean found = false;
        for (i = 0; i < curNode.size; i++) {
            if (curNode.keys[i] == key) {
                curNode.deleteLeafPointer(curNode.pointers[i]);
                found = true;
                break;
            }
        }

        if(found){
            break;
        }
        else {
            //get next node and update parents and pointer lists
            curNode = getNextNode(parents, parentPointer); //this the one , this line T-T
        }
    }
}

```

Figure 17: Remove of keys from leaf nodes (Instantiation & Removal)

Fig 17 shows the logic of the removal of leaf nodes. It requires the parameters of the specified key, a new arraylist in which stores BTreeNodes, and an array storing the subsequent pointers to key. If the key is found, it will then be successfully removed.


```

if(curNode == null){
    System.out.println("not found");
    return this;
}

if(curNode == this){
    if(curNode.size == 0){
        System.out.println("Empty tree!");
    }
    return this;
}

// if the node from deletion has lesser than minimum keys
if( curNode.size < MIN_KEYS){
    BTreeNode parentNode = parents.get(0);
    int childIndex = parentPointer.get(0);

    parentPointer.remove(0);
    // check if left sibling exist
    if(childIndex-1 >= 0){
        BTreeNode leftSibling = (BTreeNode) parentNode.pointers[childIndex-1];
        // check if left sibling has enough keys
        if(leftSibling.size>MIN_KEYS){
            //take keys from sibling
            curNode.insertKey(leftSibling.keys[leftSibling.size-1], leftSibling.pointers[leftSibling.size-1]);
            leftSibling.deleteLeafPointer(leftSibling.pointers[leftSibling.size-1]);
            //update the parent keys
            parentNode.keys[childIndex-1] = curNode.keys[0];
            return this;
        }
    }
    // check if right sibling exists
    if(childIndex+1 < parentNode.size+1){
        BTreeNode rightSibling = (BTreeNode) parentNode.pointers[childIndex+1];
        // check if right sibling has enough keys
        if(rightSibling.size>MIN_KEYS){
            curNode.insertKey(rightSibling.keys[0], rightSibling.pointers[0]);
            rightSibling.deleteLeafPointer(rightSibling.pointers[0]);
            // update the parents

```

Figure 18.1: Borrowing keys from siblings

```

// check if right sibling exists
if(childIndex+1 < parentNode.size+1){
    BTreeNode rightSibling = (BTreeNode) parentNode.pointers[childIndex+1];
    // check if right sibling has enough keys
    if(rightSibling.size>MIN_KEYS){
        curNode.insertKey(rightSibling.keys[0], rightSibling.pointers[0]);
        rightSibling.deleteLeafPointer(rightSibling.pointers[0]);
        // update the parents
        parentNode.keys[childIndex] = rightSibling.keys[0];
        return this;
    }
}
// no siblings to get keys
// attempt to merge with siblings
//check if left sibling exists
if(childIndex-1 >= 0){
    BTreeNode leftSibling = (BTreeNode) parentNode.pointers[childIndex-1];
    int j;
    for(i=leftSibling.size,j=0;j< curNode.size;j++,i++){
        leftSibling.keys[i] = curNode.keys[j];
        leftSibling.pointers[i] = curNode.pointers[j];
        leftSibling.size++;
    }
    // copy over last pointer
    leftSibling.pointers[MAX_POINTERS-1] = curNode.pointers[MAX_POINTERS-1];
    parentPointer.add(childIndex);
    BTreeNode result = removeInternal(parents, parentPointer);

    curNode.emptyAll();
    this.numOfNodes--;
    this.numOfDeleted++;
    if(result == this) return this;
    else{
        return result;
    }
}
}

```

Figure 18.2: Remove of leaf nodes (Deletion & Combination of Left Sibling)

Fig 18.1 and 18.2 shows the continuation of the removal logic of leaf nodes. If the key is not found in the array list, an error message will be returned back to the user. An error message will also return back to the user if the user decides to remove a value from the empty tree. Once a node is removed, it has to be ensured that the size of the leaf nodes will be above the minimum threshold of the minimum key. If the number of nodes fall below the threshold, it will check if any left sibling nodes exist, move the keys from the sibling nodes and update the parent keys. The nodes will then be merged to the right sibling and the current node will copy over the last pointer.

```

//merge right sibling if right sibling exists
if(childIndex+1 < parentNode.size+1){

    BTreeNode rightSibling = (BTreeNode) parentNode.pointers[childIndex+1];

    int j;
    for(i=curNode.size,j=0;j< rightSibling.size;j++,i++){
        curNode.keys[i] = rightSibling.keys[j];
        curNode.pointers[i] = rightSibling.pointers[j];
        curNode.size++;
    }
    // copy over last pointer
    curNode.pointers[MAX_POINTERS-1] = rightSibling.pointers[MAX_POINTERS-1]
    parentPointer.remove(index: 0);
    parentPointer.add(childIndex+1);
    BTreeNode result = removeInternal(parents, parentPointer);
    rightSibling.emptyALL();
    this.numOfNodes--;
    this.numOfDeleted++;
    if(result == this) return this;
    else{
        return result;
    }
}

}
return this;

```

Figure 19: Remove of leaf nodes (Combination of Right Sibling)

Fig 19 shows the continuation of the removal logic of leaf nodes. Once the left sibling node has been checked and verified, the steps taken are repeated for the right sibling node.

```

public BTreeNode removeInternal( ArrayList<BTreeNode> parentList, ArrayList<Integer> pointerList){

    // get the parent node from the makeshift stack
    BTreeNode parent = parentList.get(0);
    parentList.remove( index: 0);

    // get the pointer index in parent node
    int parentPtr = pointerList.get(0);
    pointerList.remove( index: 0);
    //if the parent is the root and only 1 key
    //lets kill the root and take the child
    if(parent.height == this.root && parent.size == 1){
        this.root--;
        return (BTreeNode) parent.pointers[0];
    }

    System.out.println("CFM delete wrong");
    System.out.println(parentPtr);
    //remove the key from parent
    //shift all keys forward
    for(int j=parentPtr; j< parent.size-1;j++){
        parent.keys[j] = parent.keys[j+1];
        parent.pointers[j] = parent.pointers[j+1];
    }

    parent.pointers[parent.size-1] = parent.pointers[parent.size];
    parent.pointers[parent.size] = null;
    parent.keys[parentPtr-1] = findMin((BTreeNode) parent.pointers[parentPtr]);

    parent.keys[parent.size-1] = -1;
    parent.pointers[parent.size] = null;
    parent.size--;
}

```

Figure 20: Updating the non leaf/root nodes (Combination of Right Sibling)

removeInternal is a recursive function that is called every time leaf nodes are merged, this is to maintain consistency. It has the same implementation as figure 17, except it is for non-leaf nodes. It returns either a new root or the current root when returning to remove function.

Experiment Results

Experiment 1 Results

For Experiment 1, our team has stored the data on the disk and fetched the respective outputs:

1. Number of Blocks
2. Size of the Database (in terms of MB)

| Block Size | Number Of Blocks | Database Size |
|------------|------------------|---------------|
| 200 B | 152903 | 29.0 MB |
| 500 B | 62960 | 29.0 MB |

```

EXPERIMENTS FOR 200MB
=====

Experiment 1:
The number of blocks: 152903
The size of database (in terms of MB): 29.0MB

```

Figure 21: Output for Experiment 1 for 200MB

```

EXPERIMENTS FOR 500MB
=====

Experiment 1:
The number of blocks: 62960
The size of database (in terms of MB): 29.0MB

```

Figure 22: Output for Experiment 1 for 500MB

Experiment 2 Results

For Experiment 2, our team has built the B+ tree based on the attribute “numVotes” by inserting the records sequentially and reported on the following statistics:

B+ Tree

1. Parameter n of B+ tree
2. Number of Nodes of B+ tree
3. Height of the B+ tree
4. Content of the root node and its 1st child node

| B+ Tree | | | | |
|------------|----|-----------------|--------|--|
| Block Size | n | Number Of Nodes | Height | Content for root node & 1st child node |
| 200 B | 16 | 131514 | 6 | Refer to Figure 23 |
| 500 B | 41 | 51957 | 5 | Refer to Figure 24 |

```
Experiment 2:
The parameter n of the B+ tree: 16
The number of nodes of the B+ tree: 131514
The height of the B+ tree, i.e the number of levels of the B+ tree: 6
The content of the root node:
Keys: [5, 7, 8, 10, 14, 18, 26, 41, 78, 198, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@3a71f4dd, MemoryPool.BTreeNode@7adf9f5f, MemoryPool.BTreeNode@85ede7b, MemoryPool.BTreeNode@5674cd4d, MemoryPool.BTreeNode@63961c42,
The content of the 1st child node:
Keys: [5, 5, 5, 5, 5, 5, 5, -1, -1, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@72ea2f77, MemoryPool.BTreeNode@33c7353a, MemoryPool.BTreeNode@681a9515, MemoryPool.BTreeNode@3af49f1c, MemoryPool.BTreeNode@19469ea2,
```

Figure 23: Experiment 2 output for block size 200MB

```
Experiment 2:  
The parameter n of the B+ tree: 41  
The number of nodes of the B+ tree: 51957  
The height of the B+ tree, i.e the number of levels of the B+ tree: 5  
The content of the root node:  
Keys: [8, 17, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1]  
Pointers: [MemoryPool.BTreeNode@46ee7fe8, MemoryPool.BTreeNode@7586e922, MemoryPool.BTreeNode@4ee285c6, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null, null]  
  
The content of the 1st child node:  
Keys: [5, 5, 5, 5, 5, 5, 5, 5, 6, 6, 6, 6, 6, 6, 6, 6, 7, 7, 7, 7, 7, 7, 8, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1]  
Pointers: [MemoryPool.BTreeNode@621be5d1, MemoryPool.BTreeNode@573fd745, MemoryPool.BTreeNode@15327b79, MemoryPool.BTreeNode@4f2410ac, MemoryPool.BTreeNode@722c41f4,
```

Figure 24: Experiment 2 output for blocksize 500MB

Experiment 3 Results

```

Experiment 3:
Accessing nodes:
Keys: [5, 7, 8, 10, 14, 18, 26, 41, 78, 198, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@a7f1f4dd, MemoryPool.BTreeNode@7ad9f9f5, MemoryPool.BTreeNode@885ede7b, MemoryPool.BTreeNode@5674cdcd, MemoryPool.BTreeNode@63961c42, MemoryPool.BTreeNode@6b5b5408, MemoryPool.BTreeNode@1be6f5c3,
Keys: [225, 259, 293, 337, 450, 646, 1821, 2181, 5251, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@1bd23cdfa, MemoryPool.BTreeNode@eedf1f4, MemoryPool.BTreeNode@7229724f, MemoryPool.BTreeNode@4c87333b, MemoryPool.BTreeNode@19d70407, MemoryPool.BTreeNode@77c8bdf, MemoryPool.BTreeNode@4eec777f,
Keys: [465, 481, 498, 508, 519, 530, 542, 553, 566, 578, 589, 640, 631, -1, -1]
Pointers: [MemoryPool.BTreeNode@a311334, MemoryPool.BTreeNode@862ab2db, MemoryPool.BTreeNode@58d75dc8, MemoryPool.BTreeNode@214c265e, MemoryPool.BTreeNode@44b139f8, MemoryPool.BTreeNode@7cca494b, MemoryPool.BTreeNode@7ba4f24f,
Keys: [499, 500, 501, 502, 503, 504, 505, 506, 507, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@c29fa335, MemoryPool.BTreeNode@1ab78a5e, MemoryPool.BTreeNode@67ef20235, MemoryPool.BTreeNode@27dc5dc6, MemoryPool.BTreeNode@94f15b24, MemoryPool.BTreeNode@15aeb7ab, MemoryPool.BTreeNode@7b23ec81,
Keys: [499, 499, 499, 499, 499, 499, 499, 499, 499, 499, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@5b48bc9f, MemoryPool.BTreeNode@6f49d9df, MemoryPool.BTreeNode@72327c9f, MemoryPool.BTreeNode@018f7f48, MemoryPool.BTreeNode@04c966a, MemoryPool.BTreeNode@2f4d37d9, MemoryPool.BTreeNode@04e50df2e,
Keys: [499, 499, 499, 499, 499, 499, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [27748856, 27875368, 28136608, 28193484, 28234566, 28515712, 29055456, 182712, 257656, 337884, null, null, null, null, null, MemoryPool.BTreeNode@12edca21]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [1658828, 776748, 778548, 876288, 908228, 945284, 1143880, 1388712, 1358968, null, null, null, null, null, null, MemoryPool.BTreeNode@34dc45dca]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [1653316, 1653312, 1702340, 1827312, 2004588, 2080532, 2219740, 2510880, 2542168, null, null, null, null, null, null, MemoryPool.BTreeNode@52cc8049]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [2953656, 3151600, 3342856, 3897312, 3988940, 4838654, 5038828, 5191348, 5345484, null, null, null, null, null, null, MemoryPool.BTreeNode@5b67f412]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [5359208, 6193856, 6494884, 6883284, 7537540, 7563228, 8069428, 8386284, 8453168, null, null, null, null, null, null, MemoryPool.BTreeNode@27973e9b]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [8453884, 8455256, 8456028, 8937456, 9071400, 9071684, 9173640, 10009228, 10045368, null, null, null, null, null, null, MemoryPool.BTreeNode@312b1daa]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [10887748, 14345684, 15219148, 15765940, 15910428, 15942940, 15967856, 16197968, 16678912, null, null, null, null, null, null, MemoryPool.BTreeNode@7530d816]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [16791280, 16965568, 17084480, 17280568, 17319172, 17363740, 17534464, 17685256, 17758968, null, null, null, null, null, null, MemoryPool.BTreeNode@27bc26da]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [1817916, 1001084, 1979496, 19824028, 20240368, 20954488, 21097516, 21158112, 21677856, null, null, null, null, null, null, MemoryPool.BTreeNode@3941a79c]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [21818512, 21868800, 21928512, 21980740, 22127410, 22239466, 22537666, 22784208, 23187080, null, null, null, null, null, null, MemoryPool.BTreeNode@8961eb77]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [23264968, 23611940, 23823484, 23985480, 24191560, 24269656, 24326828, 24587112, 26386768, null, null, null, null, null, null, MemoryPool.BTreeNode@4fca772d]
Keys: [500, 500, 500, 500, 500, 500, 500, 500, -1, -1, -1, -1, -1, -1]
Pointers: [26414080, 26684856, 26784948, 26833948, 27711256, 27741280, 28254228, 28282484, 28341284, null, null, null, null, null, null, MemoryPool.BTreeNode@9807454c]

```

Number of Nodes accessed: 19

```

data block 1
  Record 0      tt0013674      7.0      500
  Record 1      tt0024561      6.8      500
  Record 2      tt0028277      7.7      500
  Record 3      tt0041956      6.5      500
  Record 4      tt0047361      7.3      500
  Record 5      tt0047434      6.3      500
  Record 6      tt0051500      5.1      500
data block 2
  Record 7      tt0052815      6.7      500
  Record 8      tt0054298      3.7      500
  Record 9      tt0062345      4.7      500
  Record 10     tt0069064      5.8      500
  Record 11     tt0070783      5.1      500
  Record 12     tt0082275      7.1      500
  Record 13     tt0082841      5.1      500
data block 3
  Record 14     tt0085198      7.0      500
  Record 15     tt0090356      3.8      500
  Record 16     tt0099259      6.0      500
  Record 17     tt0100811      5.3      500
  Record 18     tt0106600      4.8      500
  Record 19     tt0119415      4.6      500
  Record 20     tt0120462      4.5      500
data block 4
  Record 21     tt0149695      4.3      500
  Record 22     tt0163456      8.1      500
  Record 23     tt0184456      5.7      500
  Record 24     tt0214362      8.1      500
  Record 25     tt0218000      6.3      500
  Record 26     tt0289115      5.1      500
  Record 27     tt0303487      5.9      500
data block 5
  Record 28     tt0314564      2.9      500
  Record 29     tt0327085      5.5      500
  Record 30     tt0327546      6.4      500
  Record 31     tt0398451      6.5      500
  Record 32     tt0424365      5.5      500
  Record 33     tt0450955      3.9      500
  Record 34     tt0514442      9.1      500

```

The number of data blocks accessed is 110

Average of average rating: 6.731818214329806

Figure 25: Output for Experiment 3 for 200MB

[illegible]

```
Number of Nodes accessed: 10
```


The content of data blocks accessed:

data block 1

| | | | |
|-----------|-----------|-----|-----|
| Record 0 | tt0013674 | 7.0 | 500 |
| Record 1 | tt0024561 | 6.8 | 500 |
| Record 2 | tt0028277 | 7.7 | 500 |
| Record 3 | tt0041956 | 6.5 | 500 |
| Record 4 | tt0047361 | 7.3 | 500 |
| Record 5 | tt0047434 | 6.3 | 500 |
| Record 6 | tt0051500 | 5.1 | 500 |
| Record 7 | tt0052815 | 6.7 | 500 |
| Record 8 | tt0054298 | 3.7 | 500 |
| Record 9 | tt0062345 | 4.7 | 500 |
| Record 10 | tt0069064 | 5.8 | 500 |
| Record 11 | tt0070783 | 5.1 | 500 |
| Record 12 | tt0082275 | 7.1 | 500 |
| Record 13 | tt0082841 | 5.1 | 500 |
| Record 14 | tt0085198 | 7.0 | 500 |
| Record 15 | tt0090356 | 3.8 | 500 |
| Record 16 | tt0099259 | 6.0 | 500 |

data block 2

| | | | |
|-----------|-----------|-----|-----|
| Record 17 | tt0100811 | 5.3 | 500 |
| Record 18 | tt0106600 | 4.8 | 500 |
| Record 19 | tt0119415 | 4.6 | 500 |
| Record 20 | tt0120462 | 4.5 | 500 |
| Record 21 | tt0149695 | 4.3 | 500 |
| Record 22 | tt0163456 | 8.1 | 500 |
| Record 23 | tt0184456 | 5.7 | 500 |
| Record 24 | tt0214362 | 8.1 | 500 |
| Record 25 | tt0218000 | 6.3 | 500 |
| Record 26 | tt0289115 | 5.1 | 500 |
| Record 27 | tt0303487 | 5.9 | 500 |
| Record 28 | tt0314564 | 2.9 | 500 |
| Record 29 | tt0327085 | 5.5 | 500 |
| Record 30 | tt0327546 | 6.4 | 500 |
| Record 31 | tt0398451 | 6.5 | 500 |
| Record 32 | tt0424365 | 5.5 | 500 |
| Record 33 | tt0450955 | 3.9 | 500 |

| | | | | |
|--------------|------------|-----|-----|--|
| data block 3 | | | | |
| Record 34 | tt0514442 | 9.1 | 500 | |
| Record 35 | tt0517623 | 8.1 | 500 | |
| Record 36 | tt0558715 | 8.1 | 500 | |
| Record 37 | tt0582480 | 8.3 | 500 | |
| Record 38 | tt0588114 | 8.7 | 500 | |
| Record 39 | tt0588139 | 8.1 | 500 | |
| Record 40 | tt0588187 | 8.2 | 500 | |
| Record 41 | tt0588214 | 8.6 | 500 | |
| Record 42 | tt0629708 | 8.1 | 500 | |
| Record 43 | tt0640298 | 8.0 | 500 | |
| Record 44 | tt0640308 | 8.2 | 500 | |
| Record 45 | tt0647512 | 7.2 | 500 | |
| Record 46 | tt0716912 | 7.6 | 500 | |
| Record 47 | tt0720140 | 7.1 | 500 | |
| Record 48 | tt0807689 | 8.6 | 500 | |
| Record 49 | tt11691696 | 6.4 | 500 | |
| Record 50 | tt1261908 | 8.1 | 500 | |

```
data block 4
  Record 51      tt1340802      6.0      500
  Record 52      tt1365490      6.5      500
  Record 53      tt1371697      7.9      500
  Record 54      tt1376451      5.9      500
  Record 55      tt1421383      7.6      500
  Record 56      tt1515736      7.5      500
  Record 57      tt1535989      6.4      500
  Record 58      tt1571100      5.9      500
  Record 59      tt1592534      6.6      500
  Record 60      tt1632756      7.4      500
  Record 61      tt1640740      7.7      500
  Record 62      tt1648683      7.5      500
  Record 63      tt1684558      5.3      500
  Record 64      tt1711021      6.2      500
  Record 65      tt1727519      4.1      500
  Record 66      tt1857596      7.8      500
  Record 67      tt2040651      7.8      500
```

```
data block 5
  Record 68      tt2236257      7.4      500
  Record 69      tt2247719      7.7      500
  Record 70      tt2365211      8.1      500
  Record 71      tt2662228      6.6      500
  Record 72      tt2764038      8.3      500
  Record 73      tt2813064      5.6      500
  Record 74      tt3127434      8.1      500
  Record 75      tt3203366      8.2      500
  Record 76      tt3231390      3.6      500
  Record 77      tt3263598      7.3      500
  Record 78      tt3276280      8.3      500
  Record 79      tt3398808      5.5      500
  Record 80      tt3440780      7.1      500
  Record 81      tt3592904      2.3      500
  Record 82      tt3683072      6.0      500
  Record 83      tt3916858      7.8      500
  Record 84      tt4038966      6.4      500
```

The number of data blocks accessed is 110

Average of average rating: 6.731818214329806

Figure 26: Output for Experiment 3 for 500MB

For experiment 3, we have to retrieve those movies with the “numVotes” equal to 500 and report the following statistics:

1. Number and Content of index nodes the process accesses
2. Number and Content of data blocks the process accesses
3. Average of “averageRating’s” of the records that are returned

| | | Block Size | |
|--|--|-----------------------------------|-----------------------------------|
| | | 200 B | 500 B |
| Index Nodes Accessed | Number | 19 | 10 |
| | Content of index nodes and data blocks | Refer to the outputs in Figure 25 | Refer to the outputs in Figure 26 |
| Data Blocks Accessed | Number | 110 | 110 |
| | Content of index nodes and data blocks | Refer to outputs in the Figure 25 | Refer to the outputs in Figure 26 |
| Average of “averageRating’s” of Records Returned | | 6.731818214329806 | 6.731818214329806 |

Experiment 4 Results

For experiment 4, our team has retrieved the movies with the attribute “numVotes” from 30,000 to 40,000, both inclusively and reported the following statistics:

1. Number and Content of index nodes the process accesses
2. Number and Content of data blocks the process accesses
3. Average of “averageRating’s” of the records that are returned

| | | Block Size | |
|--|---------|-------------------------------|-------------------------------|
| | | 200 B | 500 B |
| Index Nodes Accessed | Number | 91 | 38 |
| | Content | Refer to outputs in Figure 27 | Refer to outputs in Figure 28 |
| Data Blocks Accessed | Number | 942 | 930 |
| | Content | Refer to outputs in Figure 27 | Refer to outputs in Figure 28 |
| Average of “averageRating’s” of Records Returned | | 6.727911862221244 | 6.727911862221244 |

```

Experiment 4:
Accessing nodes:
Keys: [5, 7, 8, 10, 14, 18, 26, 41, 78, 198, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@3a71f4dd, MemoryPool.BTreeNode@7adf9f5f, MemoryPool.BTreeNode@85ede7b, MemoryPool.BTreeNode@5674cd4d, MemoryPool.BTreeNode@63961c42, Me

Keys: [225, 259, 293, 337, 450, 646, 1021, 2101, 5251, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@1b28cdfa, MemoryPool.BTreeNode@eed1f14, MemoryPool.BTreeNode@7229724f, MemoryPool.BTreeNode@4c873330, MemoryPool.BTreeNode@119d7047, Me

Keys: [6476, 8078, 10030, 12685, 16036, 20151, 32012, 54807, 117029, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@7e6cbb7a, MemoryPool.BTreeNode@7c3df479, MemoryPool.BTreeNode@7106e68e, MemoryPool.BTreeNode@7eda2dbb, MemoryPool.BTreeNode@6576fe71, M

Keys: [20697, 21313, 22002, 22703, 23360, 24128, 24833, 25682, 26715, 27668, 28769, 30034, 30902, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@2e817b38, MemoryPool.BTreeNode@c4437c4, MemoryPool.BTreeNode@433c675d, MemoryPool.BTreeNode@3f91beef, MemoryPool.BTreeNode@1a6c5a9e, Me

Keys: [28916, 29004, 29116, 29268, 29387, 29594, 29633, 29743, 29848, 29959, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@1d56ce6a, MemoryPool.BTreeNode@5197848c, MemoryPool.BTreeNode@17f052a3, MemoryPool.BTreeNode@2e0fa5d3, MemoryPool.BTreeNode@5010be06, Me

Keys: [29959, 29962, 29974, 29975, 29975, 29978, 29982, 29988, 29996, 30022, -1, -1, -1, -1, -1, -1]
Pointers: [20381340, 6700600, 2714084, 17457800, 22671056, 4787400, 24635712, 23120168, 20097028, 941940, null, null, null, null, null, null, MemoryPool.BTreeNode@7289

Keys: [30034, 30037, 30041, 30049, 30053, 30056, 30078, 30078, 30081, 30085, 30090, -1, -1, -1, -1, -1]
Pointers: [304340, 1862828, 22108256, 16370856, 2309200, 2196340, 4394968, 7237656, 20013256, 18977168, 1544628, null, null, null, null, null, null, MemoryPool.BTreeNode@21b

Keys: [30136, 30144, 30149, 30158, 30158, 30168, 30175, 30177, 30195, 30206, 30221, 30240, -1, -1, -1, -1]
Pointers: [11366140, 16267000, 20983456, 2256768, 4539368, 29908340, 29629456, 1638140, 17824140, 21698656, 1626968, 2373940, null, null, null, null, null, MemoryPool.BTreeN

Keys: [30246, 30247, 30248, 30254, 30259, 30262, 30275, 30319, 30326, 30333, 30341, 30354, 30361, 30370, -1, -1]
Pointers: [990000, 23826568, 28529884, 546740, 3738712, 5914368, 6406312, 14408940, 22145940, 11561084, 2196684, 6678684, 17265056, 13516828, null, null, MemoryPool.BT

Keys: [30376, 30391, 30395, 30402, 30418, 30423, 30431, 30446, 30453, 30456, 30467, 30458, 30462, -1, -1, -1]
Pointers: [7152028, 6314656, 17478800, 1063056, 23664912, 7191540, 22442856, 2259028, 29094428, 16594400, 1840428, 1801656, 2006856, null, null, null, null, MemoryPool.BTree

```

Number of Nodes accessed: 91

```

data block 1
  Record 0      tt0054167      7.7      30022
  Record 1      tt0026778      7.9      30034
  Record 2      tt0091828      5.6      30037
  Record 3      tt3361792      6.8      30041
  Record 4      tt1456941      6.2      30049
  Record 5      tt0110443      6.3      30053
  Record 6      tt0105629      5.1      30056
data block 2
  Record 7      tt0257516      5.0      30078
  Record 8      tt0489664      6.7      30078
  Record 9      tt2303687      8.7      30081
  Record 10     tt2027128      7.5      30085
  Record 11     tt0078718      7.4      30090
  Record 12     tt0857265      6.5      30136
  Record 13     tt1435513      6.7      30144
data block 3
  Record 14     tt2712740      8.1      30149
  Record 15     tt0108211      7.5      30158
  Record 16     tt0268397      6.0      30158
  Record 17     tt9053874      6.0      30168
  Record 18     tt8758202      7.5      30175
  Record 19     tt0082533      6.6      30177
  Record 20     tt1741256      7.2      30195
data block 4
  Record 21     tt3139086      7.5      30206
  Record 22     tt0082089      7.4      30221
  Record 23     tt0113253      4.9      30240
  Record 24     tt0056111      8.1      30246
  Record 25     tt4396630      7.4      30247
  Record 26     tt7668518      6.8      30248
  Record 27     tt0037382      7.8      30254
data block 5
  Record 28     tt0204993      7.4      30259
  Record 29     tt0377091      7.2      30262
  Record 30     tt0417001      5.9      30275
  Record 31     tt1174730      7.4      30319
  Record 32     tt3385524      7.2      30326
  Record 33     tt0880502      7.8      30333
  Record 34     tt0105643      2.9      30341

```

The number of data blocks accessed is 942

Average of average rating: 6.727911862221244

Figure 27: Output for Experiment 4 for 200MB

[illegible]

Number of Nodes accessed: 38

```

data block 1
    Record 0      tt0054167      7.7      30022
    Record 1      tt0026778      7.9      30034
    Record 2      tt0091828      5.6      30037
    Record 3      tt3361792      6.8      30041
    Record 4      tt1456941      6.2      30049
    Record 5      tt0110443      6.3      30053
    Record 6      tt0105629      5.1      30056
    Record 7      tt0257516      5.0      30078
    Record 8      tt0489664      6.7      30078
    Record 9      tt2303687      8.7      30081
    Record 10     tt2027128      7.5      30085
    Record 11     tt0078718      7.4      30090
    Record 12     tt0857265      6.5      30136
    Record 13     tt1435513      6.7      30144
    Record 14     tt2712740      8.1      30149
    Record 15     tt0108211      7.5      30158
    Record 16     tt0268397      6.0      30158
data block 2
    Record 17     tt9053874      6.0      30168
    Record 18     tt8758202      7.5      30175
    Record 19     tt0082533      6.6      30177
    Record 20     tt1741256      7.2      30195
    Record 21     tt3139086      7.5      30206
    Record 22     tt0082089      7.4      30221
    Record 23     tt0113253      4.9      30240
    Record 24     tt0056111      8.1      30246
    Record 25     tt4396630      7.4      30247

```


| | | | |
|--------------|-----------|-----|-------|
| Record 24 | tt0056111 | 8.1 | 30246 |
| Record 25 | tt4396630 | 7.4 | 30247 |
| Record 26 | tt7668518 | 6.8 | 30248 |
| Record 27 | tt0037382 | 7.8 | 30254 |
| Record 28 | tt0204993 | 7.4 | 30259 |
| Record 29 | tt0377091 | 7.2 | 30262 |
| Record 30 | tt0417001 | 5.9 | 30275 |
| Record 31 | tt1174730 | 7.4 | 30319 |
| Record 32 | tt3385524 | 7.2 | 30326 |
| Record 33 | tt0880502 | 7.8 | 30333 |
| data block 3 | | | |
| Record 34 | tt0105643 | 2.9 | 30341 |
| Record 35 | tt0439662 | 7.2 | 30354 |
| Record 36 | tt1629757 | 7.2 | 30361 |
| Record 37 | tt1094249 | 7.3 | 30370 |
| Record 38 | tt0482088 | 7.0 | 30376 |
| Record 39 | tt0408777 | 7.5 | 30391 |
| Record 40 | tt1672723 | 7.4 | 30395 |
| Record 41 | tt0059026 | 8.3 | 30402 |
| Record 42 | tt4283054 | 8.4 | 30418 |
| Record 43 | tt0485601 | 7.6 | 30423 |
| Record 44 | tt3544082 | 7.1 | 30431 |
| Record 45 | tt0108308 | 4.8 | 30446 |
| Record 46 | tt8201170 | 5.8 | 30453 |
| Record 47 | tt1500491 | 5.6 | 30456 |
| Record 48 | tt0090887 | 6.1 | 30457 |
| Record 49 | tt0092593 | 8.0 | 30458 |

| | | | | |
|--------------|-----------|-----------|-----|-------|
| | Record 48 | tt0090887 | 6.1 | 30457 |
| | Record 49 | tt0092593 | 8.0 | 30458 |
| | Record 50 | tt0100994 | 6.5 | 30462 |
| data block 4 | | | | |
| | Record 51 | tt0088526 | 7.9 | 30468 |
| | Record 52 | tt0105121 | 6.4 | 30492 |
| | Record 53 | tt0060665 | 7.7 | 30516 |
| | Record 54 | tt0899106 | 5.7 | 30522 |
| | Record 55 | tt0845046 | 7.0 | 30530 |
| | Record 56 | tt0098749 | 6.3 | 30540 |
| | Record 57 | tt0479968 | 4.0 | 30547 |
| | Record 58 | tt2294677 | 6.7 | 30548 |
| | Record 59 | tt1668191 | 7.4 | 30550 |
| | Record 60 | tt0043338 | 8.1 | 30552 |
| | Record 61 | tt1641384 | 8.6 | 30554 |
| | Record 62 | tt0083946 | 8.1 | 30569 |
| | Record 63 | tt5691670 | 6.5 | 30569 |
| | Record 64 | tt0096694 | 7.0 | 30571 |
| | Record 65 | tt0088814 | 6.4 | 30576 |
| | Record 66 | tt0116654 | 4.6 | 30578 |
| | Record 67 | tt1588398 | 5.3 | 30585 |
| data block 5 | | | | |
| | Record 68 | tt4131606 | 8.7 | 30605 |
| | Record 69 | tt0458413 | 6.4 | 30608 |
| | Record 70 | tt3398268 | 7.7 | 30611 |
| | Record 71 | tt0192614 | 5.6 | 30619 |

| | | | |
|--|-----------|-----|-------|
| Record 72 | tt3247714 | 5.6 | 30620 |
| Record 73 | tt0380599 | 6.8 | 30621 |
| Record 74 | tt1615918 | 4.3 | 30639 |
| Record 75 | tt0053488 | 8.6 | 30658 |
| Record 76 | tt1372686 | 6.1 | 30658 |
| Record 77 | tt2088003 | 5.4 | 30661 |
| Record 78 | tt0093148 | 6.0 | 30669 |
| Record 79 | tt0110527 | 6.5 | 30672 |
| Record 80 | tt0802999 | 8.9 | 30673 |
| Record 81 | tt0378947 | 6.4 | 30674 |
| Record 82 | tt0086687 | 7.3 | 30677 |
| Record 83 | tt3735246 | 7.2 | 30693 |
| Record 84 | tt1971352 | 6.4 | 30697 |
| The number of data blocks accessed is 930 | | | |
| Average of average rating: 6.727911862221244 | | | |

Figure 28: Output for Experiment 4 for 500MB

Experiment 5 Results

For experiment 5, our team has deleted the movies with the attribute “numVotes” equal to 1,000, updated the B+ tree accordingly, and reported the following statistics:

1. Number of times that a node is deleted during the process of updating the B+ tree
2. Number nodes of the updated B+ tree
3. Height of the updated B+ tree
4. Content of the root node and its 1st child node of the updated B+ tree

| B+ Tree | | | | | |
|------------|-----------------------------------|-----------------|--------|-------------------------------|--------------------------|
| Block Size | Number Of Times A Node Is Deleted | Number Of Nodes | Height | Content | |
| | | | | Root Node | First Child of Root Node |
| 200 B | 3 | 131511 | 5 | Refer to outputs in Figure 29 | |
| 500 B | 1 | 51956 | 4 | Refer to outputs in Figure 30 | |

```
Experiment 5:
Number of times that a node is deleted : 3
Number nodes of the updated B+ tree: 131511
Height of the updated B+ tree: 5
content of the root node:
Keys: [5, 7, 8, 10, 14, 18, 26, 41, 78, 198, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@3a71f4dd, MemoryPool.BTreeNode@7adf9f5f, MemoryPool.BTreeNode@85ede7b, MemoryPool.BTreeNode@5674cd4d, MemoryPool.BTreeNode@63961c42,
MemoryPool.BTreeNode@72ea2f77, MemoryPool.BTreeNode@33c7353a, MemoryPool.BTreeNode@681a9515, MemoryPool.BTreeNode@3af49f1c, MemoryPool.BTreeNode@19469ea2]
Keys: [5, 5, 5, 5, 5, 5, 5, 5, -1, -1, -1, -1, -1, -1, -1, -1]
Pointers: [MemoryPool.BTreeNode@72ea2f77, MemoryPool.BTreeNode@33c7353a, MemoryPool.BTreeNode@681a9515, MemoryPool.BTreeNode@3af49f1c, MemoryPool.BTreeNode@19469ea2]

Please choose an option to continue:

[1] Repeat Experiments 1-5 with new block size of 500MB
[2] Exit
```

Figure 29: Output for Experiment 5 for 200MB

Figure 30: Output for Experiment 5 for 500MB